



# Applied Containerization for Machine Learning in HPC

# Applied Containerization for Machine Learning in HPC

Date: October 21, 2025

Instructor: Mohal Khandelwal

- Website: [www.rc.colorado.edu/rc](http://www.rc.colorado.edu/rc)
- Documentation: <https://curc.readthedocs.io>
- Helpdesk: [rc-help@colorado.edu](mailto:rc-help@colorado.edu)
- Survey: <http://tinyurl.com/curc-survey18>



**Slides**

[https://github.com/ResearchComputing/applied\\_containerization\\_for\\_ml\\_short\\_course](https://github.com/ResearchComputing/applied_containerization_for_ml_short_course)



# Meet the User Support Team



Layla  
Freeborn



Brandon  
Reyes



Andy  
Monaghan



Michael  
Schneider



John  
Reiland



Dylan  
Gottlieb



Mohal  
Khandelwal



Ragan  
Lee

# Why Containers for ML in HPC?

- **Dependency Management**

- Package complex ML software stacks with all dependencies.
- Resolve version conflicts and avoid "dependency hell."

- **Portability**

- Run the same environment across different HPC systems, local machines, or cloud platforms.
- "Build once, run anywhere."

- **Scalability**

- Easily scale workloads from a laptop to thousands of nodes.
- Integrate seamlessly with schedulers like SLURM for large ML jobs.

# Why Containers for ML in HPC? (cont.)

- **Reproducibility**

- Capture the exact software environment for consistent, repeatable results.

- **Performance**

- Near-native speed with minimal overhead; direct access to GPUs and high-speed interconnects.
- Optimized images can boost job throughput and resource utilization.

# What is a container?

- An isolated, encapsulated user-space instance.
- Runs on a shared OS kernel but has its own:
  - Filesystem
  - Processes
  - Network interfaces (can be configured)
- Container Image: A self-contained read-only file (or files) used to run the packaged application
- Container: A running instance of a container image



# Containers vs. Virtual Machines (VMs)

Feature	Virtual Machine (VMs)	Containers
Isolation	Full (hardware/emulated)	Process/user-space
Guest OS	Full OS per VS	Shares host OS kernel
Setup	Slow (minutes)	Fast (seconds)
Resource Use	Heavy (more RAM/CPU)	Lightweight (minimal overhead)
Portability	Hypervisor-dependent	Build once, run anywhere
Use case	Legacy apps, OS-level isolation	ML, microservices, HPC, CI/CD



# Key Components

- **Container Images**

- Read-only templates used to create containers.
- Contain application code, libraries, dependencies, and metadata.
- Examples: Docker Hub images, SIF files (Apptainer).

- **Runtime Environments**

- Software that runs containers (e.g., Apptainer, Docker).
- Manages container lifecycle, isolation, and resource allocation.



# Key Components (cont.)

- **Mount Points for Data (Bind Mounts)**
  - Mechanism to make host directories/files accessible inside the container.
  - Essential for accessing datasets, scripts, and output directories.
- **Resource Allocation**
  - Containers share host resources (CPU, memory, GPUs).
  - HPC schedulers (like SLURM) manage resource allocation for containerized jobs.

# Apptainer (formerly Singularity)

- Apptainer is the direct successor to Singularity.
- Requires a Linux system
- Runs without requiring root access
- Single-file SIF format is easy to transport and share with others
- Can convert existing Docker images to Apptainer images
- Apptainer comes pre-installed on all Alpine compute nodes, so no need to load any specific software modules!



# Using Apptainer on HPC

- View list of Apptainer commands.

```
$ apptainer --help
```

- Look at CURC's collection of pre-build containers.

```
$ echo $CURC_CONTAINER_DIR
```

```
$ ls $CURC_CONTAINER_DIR
```

- A Singularity Definition File (or “def file” for short) is a set of blueprints explaining how to build a custom container.
  - It includes specifics about the base OS, software to install, environment variables and other metadata.
  - [https://apptainer.org/docs/user/1.0/definition\\_files.html](https://apptainer.org/docs/user/1.0/definition_files.html)

# Using Apptainer on HPC

- By default, the cache directory for Apptainer builds is **/scratch/alpine/\$USER**

```
$ echo $APPTAINER_CACHEDIR
```

- By default, only /home/\$USER is available within any given container.
- To bind any additional folders/files to your container, use the -B flag in your apptainer run, exec, and shell commands:

```
$ apptainer run -B /source/directory:/target/directory sample-image.sif
```

```
$ apptainer run -B /projects/$USER,/pl/active,/scratch/alpine/$USER sample-image.sif
```



# Apptainer Commands

Other useful Apptainer commands:

<code>apptainer inspect</code>	<code>#See labels/environment vars, run scripts</code>
<code>apptainer pull</code>	<code>#pull an image from hub</code>
<code>apptainer exec</code>	<code>#Execute a command to your container</code>
<code>apptainer run</code>	<code>#Run your image as an executable</code>
<code>apptainer build</code>	<code>#Build a container</code>
<code>apptainer shell</code>	<code>#Access the command line of your container</code>

# Using images from a pre-built container

- You can fetch container images from container registries such as:
  - [Docker Hub](#)
  - [NVIDIA NGC Catalog](#)
- Let's walk through an example of using pre-built Tensorflow container to train a classification model

# TensorFlow Example with Apptainer

## 1. Export paths

```
$ export IMAGES=/projects/$USER/containers/  
$ export WORKDIR=/projects/$USER/containers/ml_work  
$ mkdir -p $IMAGES $WORKDIR && cd $WORKDIR
```

## 2. Pull TensorFlow container (GPU version)

```
$ apptainer pull $IMAGES/tensorflow-2.15.0-cuda12.8.sif  
docker://tensorflow/tensorflow:2.15.0-gpu
```

## 3. Inspect container metadata

```
$ apptainer inspect $IMAGES/tensorflow-2.15.0-cuda12.8.sif
```

# TensorFlow Example (cont.)

- Bind mount your script directory.
- Create `mnist_tf.py` that builds and trains a simple image classification model
- `--nv`: Enables NVIDIA GPU access.

## 4. Bind paths

```
$ aptainer exec -B $WORKDIR:/work,$IMAGES:/images $IMAGES/tensorflow-2.15.0-cuda12.8.sif ls /work
```

## 5. Run the script

```
$ aptainer exec --nv -B $WORKDIR:$WORKDIR $IMAGES/ tensorflow-2.15.0-cuda12.8.sif python3 $WORKDIR/mnist_tf.py
```



# Example python script

```
cat >mnist_tf.py <<'PY'
import tensorflow as tf
mnist = tf.keras.datasets.mnist
(x_train, y_train), _ = mnist.load_data()
x_train = x_train / 255.0
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28,28)),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax')
])
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
metrics=['accuracy'])
model.fit(x_train, y_train, epochs=1)
print("Done training.")
PY
```



# Access container interactively

- Shell into it

```
$ aptainer shell --nv $IMAGES/tensorflow-2.15.0-cuda12.8.sif
```

- Then run

```
$ python3
>>> import tensorflow as tf
>>> tf.__version__
>>> tf.config.list_physical_devices('GPU')
```

If it returns a list (e.g. [PhysicalDevice(name='/physical\_device:GPU:0', ...)]), then TensorFlow successfully detected the GPU inside the container!

# SLURM Integration

```
#!/bin/bash
#SBATCH --gres=gpu:1
#SBATCH --partition=aa100
#SBATCH --ntasks=4
#SBATCH --nodes=1
#SBATCH --qos=normal
#SBATCH --time=1:00:00
#SBATCH --job-name=tf-classify
#SBATCH --output=tf-classify.%j.out
#SBATCH --mail-type=ALL
#SBATCH --mail-user=<your email>

export IMAGES=/projects/$USER/containers/sif
export WORKDIR=/projects/$USER/containers/ml_work
mkdir -p $IMAGES $WORKDIR && cd $WORKDIR

apptainer exec --nv -B $WORKDIR:$WORKDIR $IMAGES/tensorflow-2.20.0.sif
python3 $WORKDIR/mnist_tf.py
```



# Troubleshooting

- Check GPU access:
  - `apptainer exec --nv $IMAGES/tensorflow-2.15.0-cuda12.8.sif python3 -c "import tensorflow as tf; print(tf.config.list_physical_devices('GPU'))"`
- File system binding:
  - Ensure your data/scripts are accessible inside the container.
- Environment variables:
  - Pass with `--env` or set inside the container.
- Cache management:
  - `apptainer cache list` and `apptainer cache clean`
- Resource monitoring:
  - Use `htop`, `nvidia-smi`, or `squeue` to monitor jobs.





# Pytorch Example with Apptainer

## 1. Export paths

```
$ export IMAGES=/projects/$USER/containers/  
$ export WORKDIR=/projects/$USER/containers/ml_work  
$ mkdir -p $IMAGES $WORKDIR && cd $WORKDIR
```

## 2. Pull TensorFlow container (GPU version)

```
$ apptainer pull $IMAGES/pytorch-2.9.0-cuda12.8.sif docker://pytorch/pytorch:2.9.0-cuda12.8-cudnn9-runtime
```

## 3. Inspect container metadata

```
$ apptainer inspect $IMAGES/ pytorch-2.9.0-cuda12.8.sif
```

# Pytorch Example (cont.)

- Bind mount your script directory.
- Create polyfit.py that builds and trains a small neural network

## 4. Bind paths

```
$ aptainer exec -B $WORKDIR:/work,$IMAGES:/images $IMAGES/pytorch-2.9.0-cuda12.8.sif ls /work
```

## 5. Run the script

```
$ aptainer exec --bind $WORKDIR:$WORKDIR $IMAGES/pytorch-2.9.0-cuda12.8.sif python3 $WORKDIR/polyfit.py
```

# Sample python script

```
cat >polyfit.py <<'PY'
import torch
x = torch.linspace(-1, 1, 100).unsqueeze(1)
y = 3 * x ** 2 + 2 * x + 1 + 0.1 * torch.randn(x.size())
model = torch.nn.Sequential(torch.nn.Linear(1, 10), torch.nn.ReLU(),
torch.nn.Linear(10, 1))
loss_fn = torch.nn.MSELoss()
optimizer = torch.optim.Adam(model.parameters())
for epoch in range(100): optimizer.zero_grad(); loss_fn(model(x),
y).backward(); optimizer.step()
print("Done training.")
PY
```



# Best Practices: Container Management

- **Version Control & Tagging**

- Tag containers with specific versions (e.g., myimage:1.0.0, myimage:latest-cuda11.8).
- Store definition files (.def, Dockerfiles) in version control (Git).

- **Data Management**

- Use bind mounts for large datasets to avoid including them in images.
- Keep images small and focused on software environment.

- **Resource Allocation**

- Match container resource needs to SLURM (or other scheduler) requests.
- Avoid over-subscribing resources.



**Any Questions**



# Thank you!

## Survey and feedback

<http://tinyurl.com/curc-survey18>

